

# Cassandra & HBase Comparison

郭鹏 盛大在线  
新浪微博：@逖靖寒

SACC2011

# Agenda

---

- ▣ Function
- ▣ Implement
- ▣ What's Next

# Agenda

---

- ▣ Function
- ▣ Implement
- ▣ What's Next

# Data Model



## Standard Column Family

Keyspace	Row	Column Family	Column Name	Column Value	Timestamp
KS-1	Row-1	cfName-1	colName-1	colValue-1	12345678
KS-1	Row-1	cfName-1	colName-2	colValue-2	22345638
KS-1	Row-1	cfName-2	colName-1	colValue-1	12445678
KS-1	Row-2	cfName-1	colName-1	colValue-1	12345678
KS-2	Row-1	cfName-1	colName-1	colValue-1	12345678

SACCC2011

# Data Model



## Super Column Family

Keyspace	Row	Column Family	Super Column Name	Column Name	Column Value	Timestamp
KS-1	Row-1	cfName-1	sColName-1	colName-1	colValue-1	12345678
KS-1	Row-1	cfName-1	sColName-1	colName-2	colValue-2	22345638
KS-1	Row-1	cfName-2	sColName-2	colName-1	colValue-1	12445678

# Data Model



Table	Row	Column Family	Column Name	Column Value	Timestamp	Version
Tab-1	Row-1	cfName-1	colName-1	colValue-1	12345678	V-1
Tab-1	Row-1	cfName-1	colName-2	colValue-2	22345638	V-2
Tab-1	Row-1	cfName-1	colName-2	colValue-2	12445678	V-1
Tab-1	Row-2	cfName-1	colName-1	colValue-1	12345678	V-1
Tab-2	Row-1	cfName-1	colName-1	colValue-1	12345678	V-1

SACC2011

# Basic API



- CRUD
- ConsistencyLevel
  - ONE
  - QUORUM
  - ALL
- Range Query :
  - Row
  - Token

# Advanced API



- Authorization & Authentication
- Secondary Index
- Online Schema Update
- Counter
- MapReduce
- CQL(Cassandra Query Language)

```
USE WebSiteKS;
SELECT FROM Standard1 WHERE KEY = "k";
UPDATE Standard1 WITH ROWC("k", COLC("c", "hello!"));
```



# Basic API



- ▣ CRUD
- ▣ Scan:
  - ▣ Row

# Advanced API



- ▣ Lock
  - ▣ Put
  - ▣ checkAndPut
  - ▣ Delete
  - ▣ checkAndDelete
- ▣ Limited Online Schema Update
- ▣ Counter
- ▣ MapReduce
- ▣ Filter

```
byte[] startRow = Bytes.toBytes(rowPrefix);  
  
Scan scan = new Scan();  
  
Filter filter = new PrefixFilter(startRow);  
scan.setFilter(filter);  
scan.setStartRow(startRow);
```

SACCC2011

# Agenda

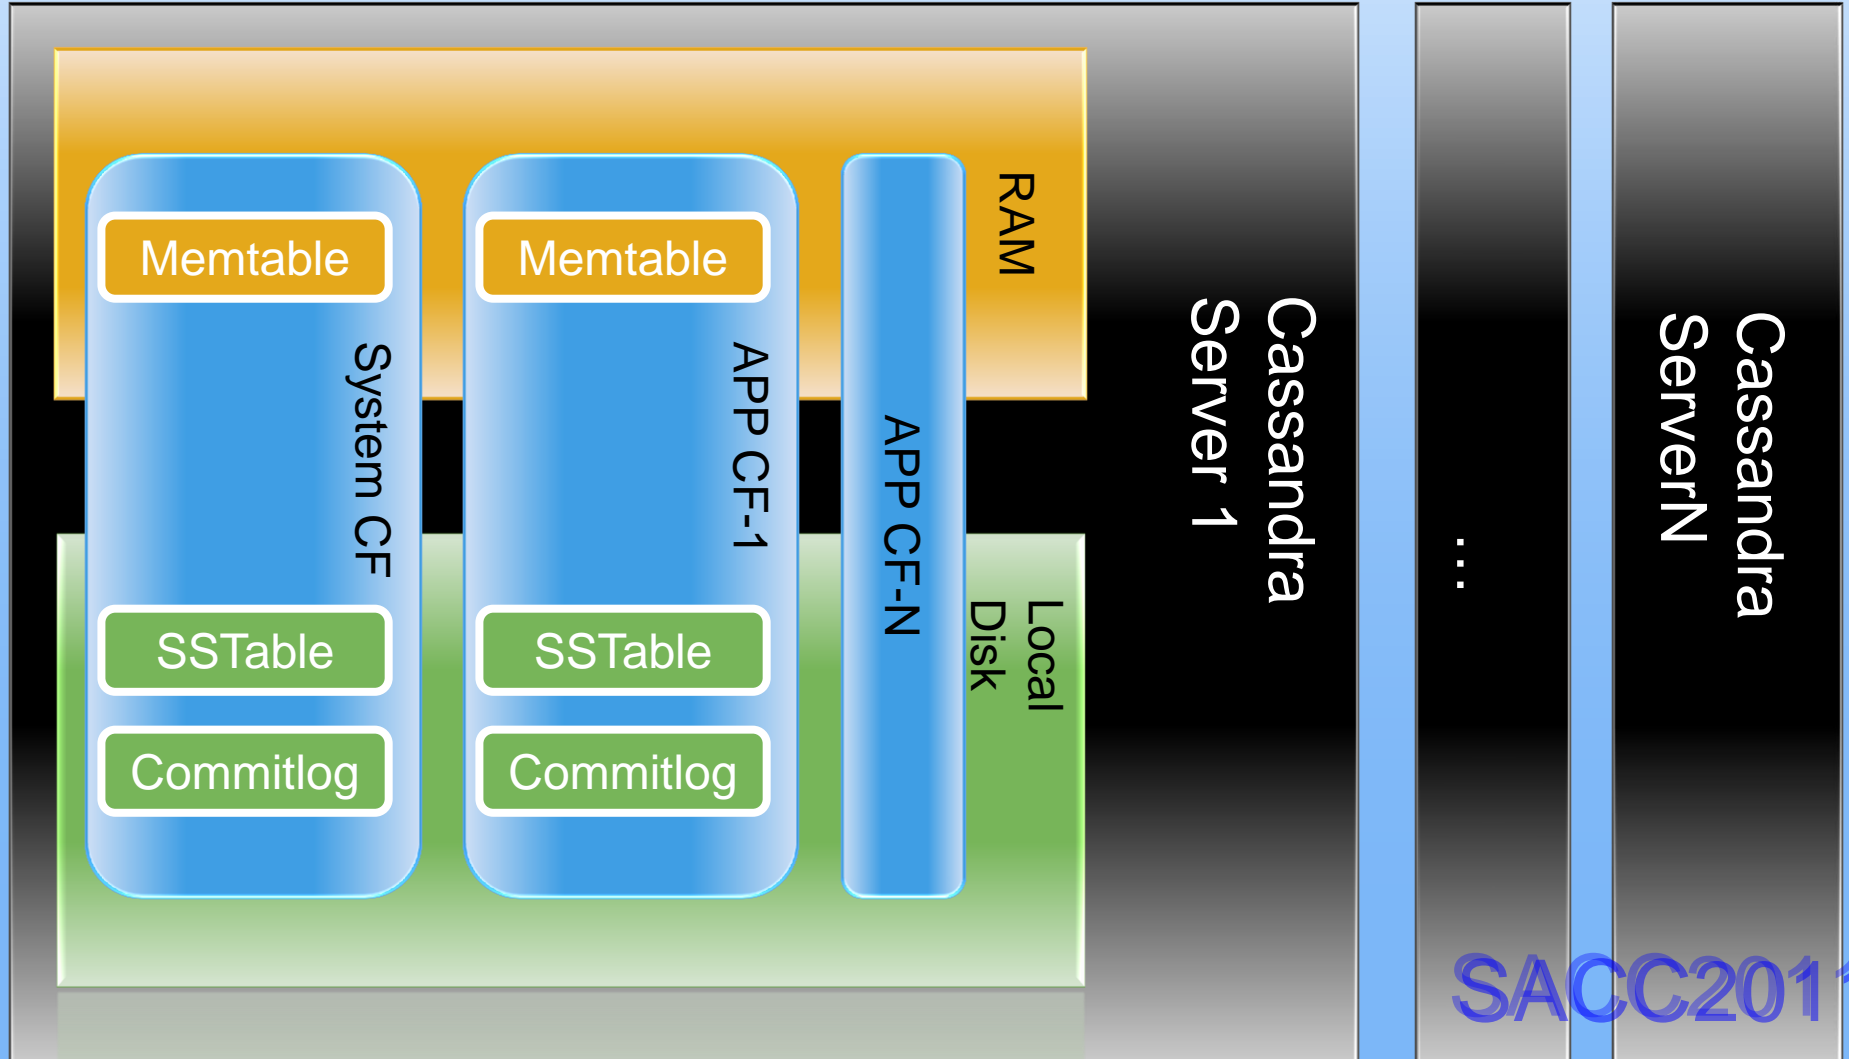
---

- ▣ Function
- ▣ Implement
- ▣ What's Next

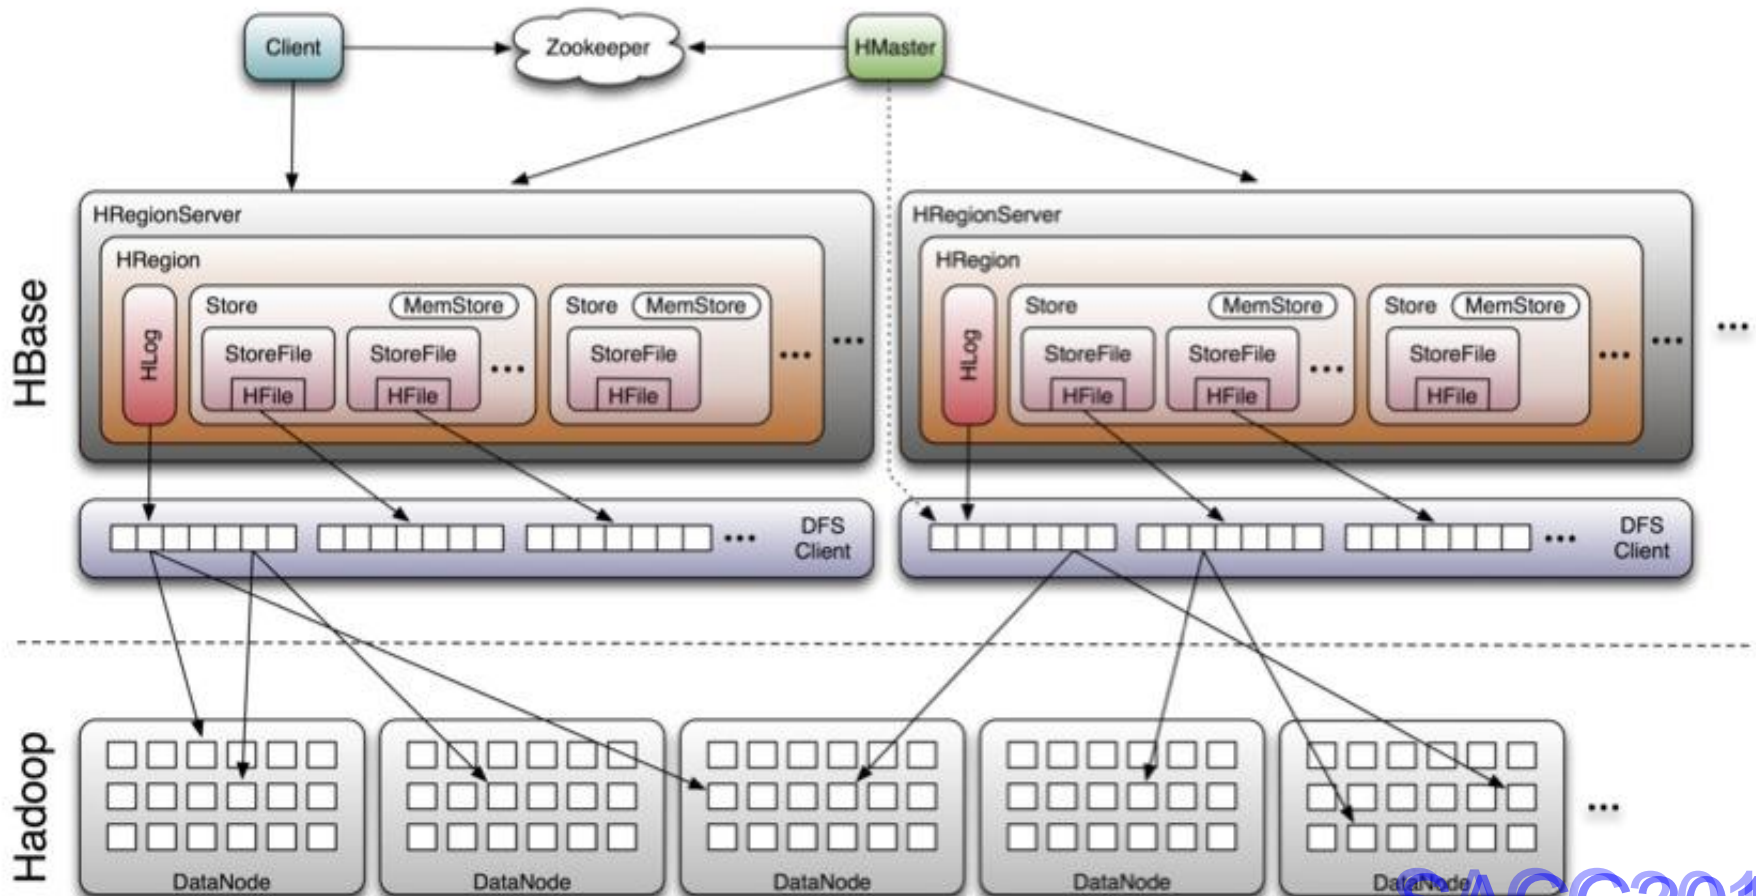
# System Architecture



**Cassandra**

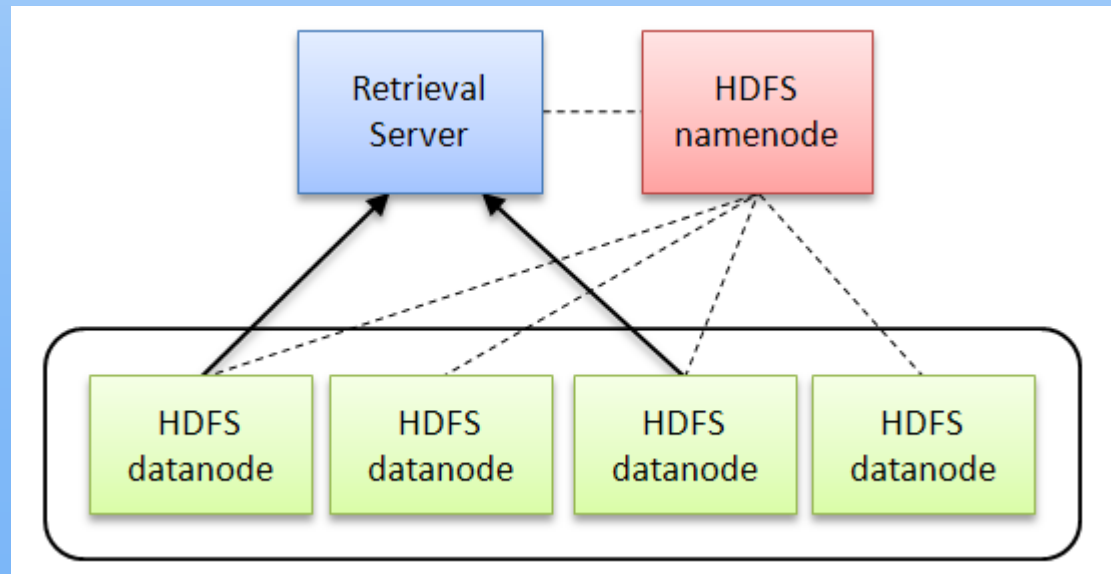


# System Architecture



SACC2011

# File System



SACC2011


# Consistency



Server No.	RowKey List
Cassandra-01	key1, <b>key2,key3,key4</b>
Cassandra-02	<b>key2,key3,key4</b> ,key5
Cassandra-03	<b>key3,key4</b> ,key5,key6

Server No.	RowKey List
RegionServer-01	key1,key2,key3,key4
RegionServer-02	key5,key6,key7,key8
RegionServer-03	key9,key10,key11,key12

# Location



Address	Status	State	Load	Token
				1167911881211208861175136640
172.16.197.113	Up	Normal	17.56 KB	7653251155202617177471488196
172.16.197.114	Up	Normal	17.83 KB	7653251155202617177471488196
172.16.197.112	Up	Normal	22.5 KB	1167911881211208861175136640

SAGC2011

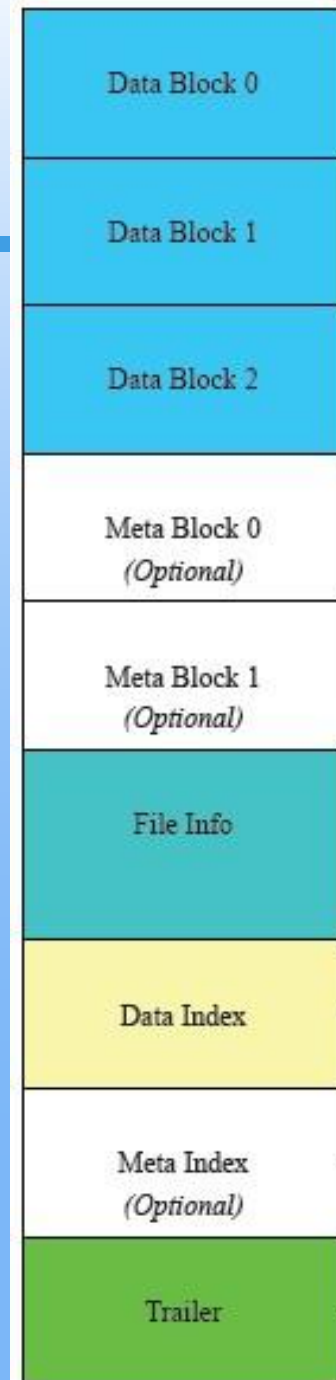



# Location



```
hbase(main):004:0> scan '.META.', {LIMIT => 1}
ROW                                COLUMN+CELL
BE,,1312597904196.00716b4aac6b3562c8fbb15afd680f94. column=info:regioninfo, timestamp=1312597904369, value=REGION
4aac6b3562c8fbb15afd680f94. => 'BE,,1312597904196.00716b4aac6b3562c8fbb15afd680f94.', STARTKEY => 'BE,,1312597904196.00716b4aac6b3562c8fbb15afd680f94.', ENDKEY => '', ENCODED => 00716b4aac6b3562c8fbb15afd680f94,
{{NAME => 'BE', FAMILIES => [{NAME => 'course', BLOOMFILTER => REPLICATION_SCOPE => '0', COMPRESSION => 'NONE', VERSIONS => '3', TTL => '2147483647', BLOCKSIZE => '65536', IN_MEMORY => 'false', BLOCKCACHE => 'true'}], {NAME => 'student', BLOOMFILTER => 'NONE', REPLICATION_SCOPE => '0', COMPRESSION => 'NONE', VERSIONS => '3', TTL => '2147483647', BLOCKSIZE => '65536', IN_MEMORY => 'false', BLOCKCACHE => 'true'}], {NAME => 'teacher', BLOOMFILTER => 'NONE', REPLICATION_SCOPE => '0', COMPRESSION => 'NONE', VERSIONS => '3', TTL => '2147483647', BLOCKSIZE => '65536', IN_MEMORY => 'false', BLOCKCACHE => 'true'}}}}
BE,,1312597904196.00716b4aac6b3562c8fbb15afd680f94. column=info:server, timestamp=1314342637415, value=hadoop02.cl
4aac6b3562c8fbb15afd680f94. :60020
BE,,1312597904196.00716b4aac6b3562c8fbb15afd680f94. column=info:serverstartcode, timestamp=1314342637415, value=1314342637415
```

# Storage File



SACC2011


# Cache



Cache	Value
RowCache	KeyIndex
KeyCache	Key/Value



Cache	Value
BlockCache	Data Block



# BulkLoader



BinaryMemtable

HRegionInterface.bulkLoadHFile

- TableName
- RegionName
- CFName

nodetool cleanup

sstableloader

LoadIncrementalHFiles

- HFilePath
- TableName

SACC2011

# LoadBalance



nodetool move token

balancer

move regionName serverName

split regionName

SACC2011

# Agenda

---

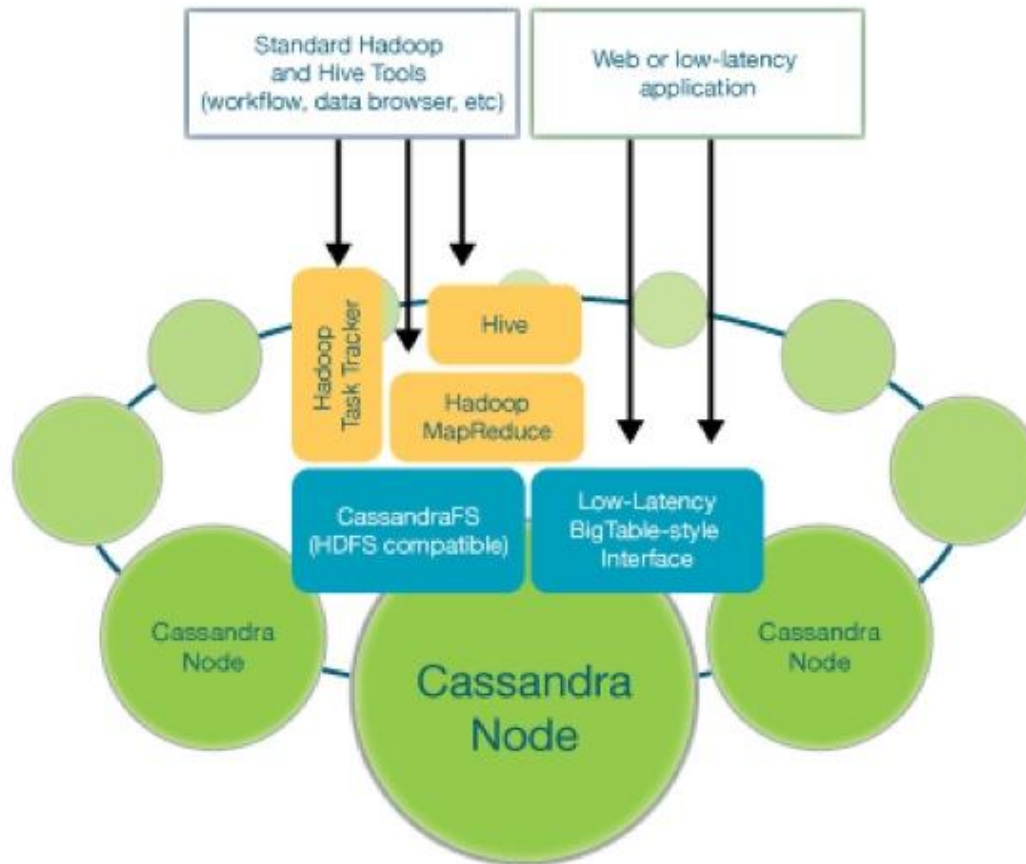
- ▣ Function
- ▣ Implement
- ▣ What's Next

- Enhanced Hadoop and Hive distribution that utilizes Cassandra for all of its core services.
- Integrated with Hadoop MapReduce and Hive
- Providing an HDFS compatible storage layer powered by Cassandra

# Application Stack



**Cassandra**



- Dynamo-based fault-tolerant replication
- Distributed peer architecture – no name node or other central point of failure

SACC2011



# Deploy



**Cassandra**




SACC2011

# Coprocessor




- ▣ MasterObserver
- ▣ RegionObserver
- ▣ EndPoint

# MasterObserver



SACCC2011

# RegionObserver

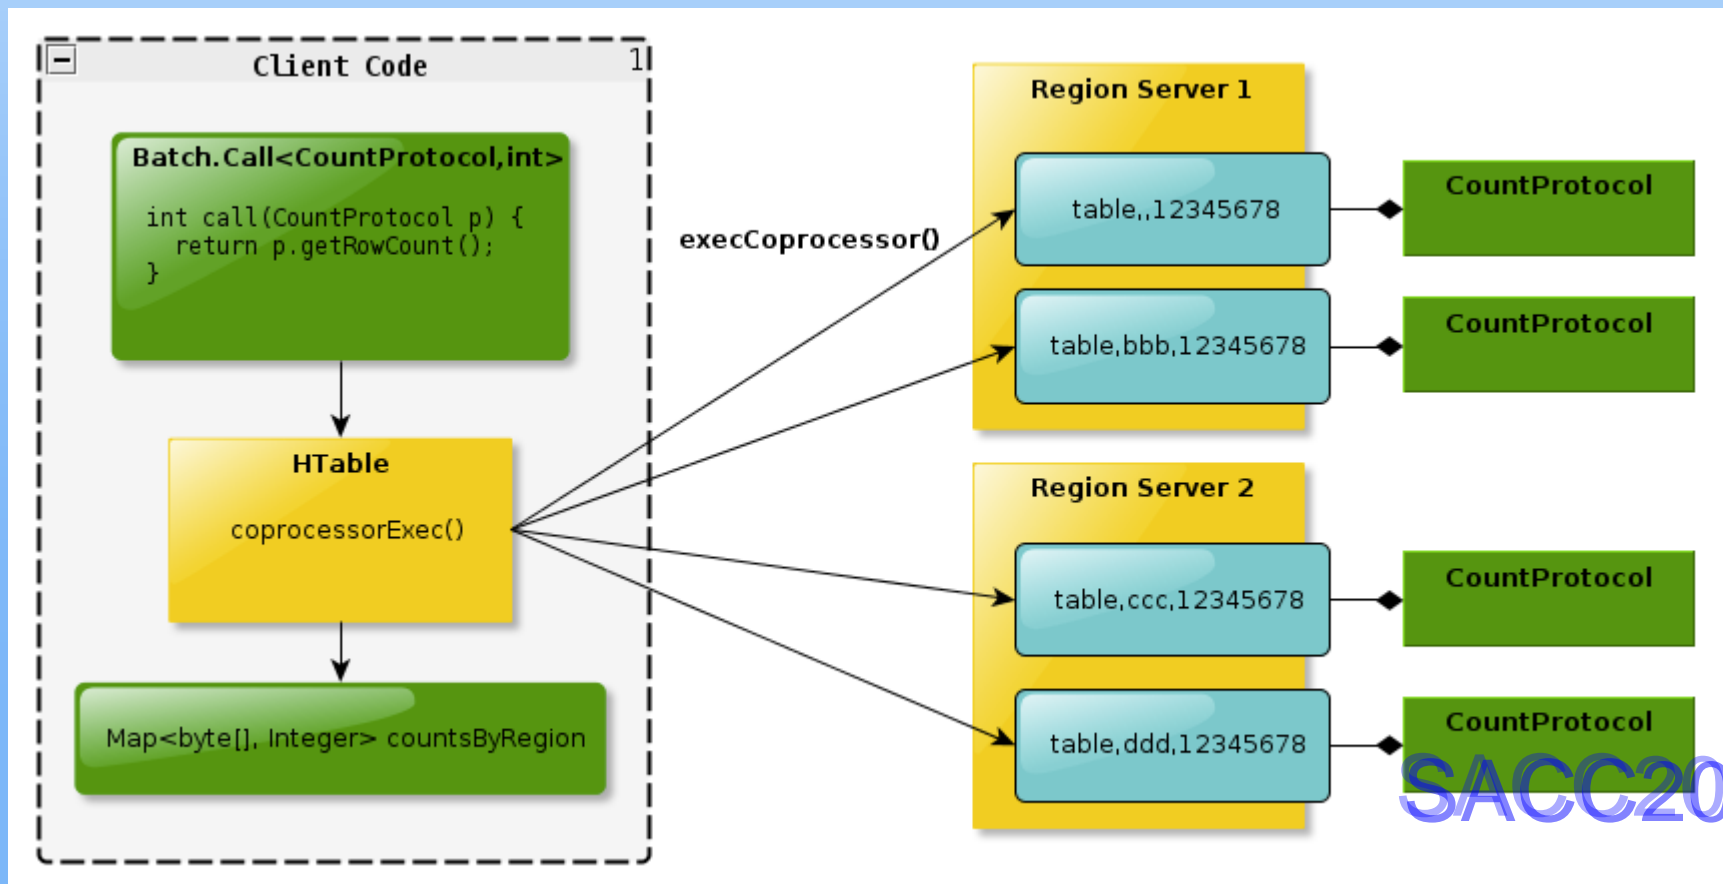


SACCC2011

# EndPoint



Define your own extensions to HBase RPC transactions exchanged between clients and the region servers.



# Applications



- HBase access control: provides basic access control for HBase
- Column aggregate: SQL-like sum(), avg(), max(), min(), etc.
- Region level indexing
- HBase table, region access statistic

# 使用场景

- ▣ Key Value
- ▣ Key Scan
- ▣ Column动态添加
- ▣ 排序
- ▣ 二级索引

# Cassandra适用场景

- ▣ 从零开始
- ▣ 快速开发
- ▣ 快速部署
- ▣ 简单的安全认证



# HBase适用场景

- ▣ 已有Hadoop集群
- ▣ 要求数据一致性
- ▣ 高效的数据平衡
- ▣ 服务端代码执行

# Thanks!

盛大在线  
WWW.SDO.COM

SACC2011